

Simple Housing Dataset

Create a regression model that predicts the price of boston house

```
In [145]: !wget https://storage.googleapis.com/nicksdemobucket/housing-data.csv

--2019-03-26 01:02:40--  https://storage.googleapis.com/nicksdemobucket/housing-data.csv
Resolving storage.googleapis.com (storage.googleapis.com)... 74.125.195.128,
2607:f8b0:400e:c08::80
Connecting to storage.googleapis.com (storage.googleapis.com)|74.125.195.128|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 816 [application/octet-stream]
Saving to: 'housing-data.csv.2'

housing-data.csv.2  100%[=====>]      816  --.-KB/s    in 0s

2019-03-26 01:02:40 (10.6 MB/s) - 'housing-data.csv.2' saved [816/816]
```

```
In [146]: # Import necessary libraries and get the housing data
import pandas as pd
import numpy as np
df = pd.read_csv('housing-data.csv')
df.head()
```

Out[146]:

	sqft	bdrms	age	price
0	2104	3	70	399900
1	1600	3	28	329900
2	2400	3	44	369000
3	1416	2	49	232000
4	3000	4	75	539900

```
In [147]: # Featurize and Label the dataset
features = df.drop(['price'], axis=1).values
labels = df[['price']].values
print(features.shape, labels.shape)

(47, 3) (47, 1)
```

```
In [0]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
```

```
In [0]: # Build the NN regression model
model = Sequential()
model.add(Dense(25, activation='relu', input_shape=(3,)))
model.add(Dense(25, activation='relu'))
model.add(Dense(1))

model.compile(optimizer='adam',
              loss='mean_squared_error',
              metrics=['mean_absolute_error'])

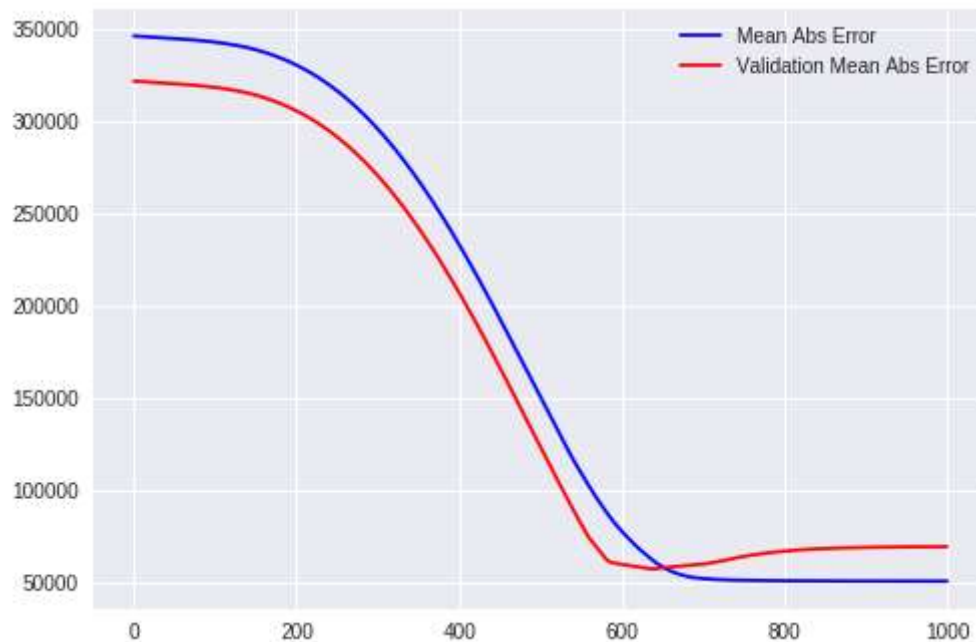
h = model.fit(features, labels, epochs=1000, validation_split=0.2,
              batch_size=130, verbose=0)
```

```
In [150]: h.history.keys()
```

```
Out[150]: dict_keys(['loss', 'mean_absolute_error', 'val_loss', 'val_mean_absolute_error'])
```

```
In [151]: # Plot the history
import matplotlib.pyplot as plt
plt.plot(h.history['mean_absolute_error'], color='b',
         label='Mean Abs Error')
plt.plot(h.history['val_mean_absolute_error'], color='r',
         label='Validation Mean Abs Error')
plt.legend()
```

```
Out[151]: <matplotlib.legend.Legend at 0x7f108424ac18>
```



Binary Classification

Create a binary classifier for the titanic dataset, will person x survive?

In [152]: !wget https://storage.googleapis.com/nicksdemobucket/titanic-train.csv

```
--2019-03-26 01:02:52-- https://storage.googleapis.com/nicksdemobucket/titanic-train.csv
Resolving storage.googleapis.com (storage.googleapis.com)... 74.125.195.128,
2607:f8b0:400e:c09::80
Connecting to storage.googleapis.com (storage.googleapis.com)|74.125.195.128|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 61194 (60K) [application/octet-stream]
Saving to: 'titanic-train.csv.2'
```

titanic-train.csv.2 100%[=====>] 59.76K --.-KB/s in 0.001s

2019-03-26 01:02:52 (88.3 MB/s) - 'titanic-train.csv.2' saved [61194/61194]

In [153]: *# Get the titanic data*
df = pd.read_csv('titanic-train.csv')
df.head()

Out[153]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	F
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.25
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.28
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.92
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.10
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.05

```
In [0]: # Remove features that seem unnecessary off the bat
df = df.drop(['PassengerId', 'Name', 'Ticket'], axis=1)
```


```
In [155]: df.head()
```

```
Out[155]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Cabin	Embarked
0	0	3	male	22.0	1	0	7.2500	NaN	S
1	1	1	female	38.0	1	0	71.2833	C85	C
2	1	3	female	26.0	0	0	7.9250	NaN	S
3	1	1	female	35.0	1	0	53.1000	C123	S
4	0	3	male	35.0	0	0	8.0500	NaN	S

```
In [156]: # Visualize the null data in the dataset
import seaborn as sns
sns.heatmap(df.isnull(), cbar=False, yticklabels=False)
```

```
Out[156]: <matplotlib.axes._subplots.AxesSubplot at 0x7f10841aae80>
```




```
In [0]: """
Since there are only a few missing rows for Embarked we will
drop the missing rows entirely
"""
df = df.dropna(subset=['Embarked'])
```

```
In [0]: """  
Remove the column "Cabin" from features as a lot of the data is missing and  
we already have a proxy for socio-economic status (Pclass)  
"""  
df = df.drop(['Cabin'], axis=1)
```

```
In [159]: # Investigate a good proxy for age - I suggest PClass  
plt.figure(figsize=(10,7))  
sns.boxplot(x='Pclass', y='Age', data=df)
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:454: FutureWarn  
ing: remove_na is deprecated and is a private function. Do not use.  
    box_data = remove_na(group_data)
```

```
Out[159]: <matplotlib.axes._subplots.AxesSubplot at 0x7f10841dbd30>
```




```
In [0]: """
Create a function to impute the age based on the average age of each PClass
"""
def impute_age(cols):
    Age = cols[0]
    Pclass = cols[1]

    if pd.isnull(Age):
        if Pclass==1:
            return 37
        elif Pclass==2:
            return 29
        else:
            return 24
    else:
        return Age
```

```
In [0]: # Apply impute_age to the dataset to fill in the missing age data
df['Age'] = df[['Age', 'Pclass']].apply(impute_age, axis=1)
```

```
In [162]: # Show that all missing data is gone
sns.heatmap(df.isnull(), cbar=False, yticklabels=False)
```

```
Out[162]: <matplotlib.axes._subplots.AxesSubplot at 0x7f10840ec358>
```



```
In [0]: # One-Hot-Encode the Sex and Embarked columns
df = pd.concat([df, pd.get_dummies(df['Sex'], prefix='Sex', drop_first=True)],
axis=1).drop(['Sex'],axis=1)
df = pd.concat([df, pd.get_dummies(df['Embarked'], prefix='Embarked', drop_fir
st=True)], axis=1).drop(['Embarked'],axis=1)
```

In [164]: `df.head()`

Out[164]:

	Survived	Pclass	Age	SibSp	Parch	Fare	Sex_male	Embarked_Q	Embarked_S
0	0	3	22.0	1	0	7.2500	1	0	1
1	1	1	38.0	1	0	71.2833	0	0	0
2	1	3	26.0	0	0	7.9250	0	0	1
3	1	1	35.0	1	0	53.1000	0	0	1
4	0	3	35.0	0	0	8.0500	1	0	1

In [165]: *# Now split the data into features and labels*
`features = df.drop(['Survived'], axis=1)`
`labels = df[['Survived']].values`
`print(features.shape, labels.shape)`

(889, 8) (889, 1)

In [0]: *# Now that the features are cleaned up - we can train the binary classifier:*

```
model = Sequential()

model.add(Dense(20, activation='relu', input_shape=(8,)))
model.add(Dense(20, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy'])


h = model.fit(features, labels, epochs=500, validation_split=0.3, verbose=0)
```

In [167]: `h.history.keys()`

Out[167]: `dict_keys(['loss', 'acc', 'val_loss', 'val_acc'])`

```
In [168]: # Plot the accuracy during training
plt.plot(h.history['acc'], color='b', label='Accuracy')
plt.plot(h.history['val_acc'], color='r', label='Validation Accuracy')
plt.legend()
```

Out[168]: <matplotlib.legend.Legend at 0x7f1083a39eb8>



MNIST

See how well a fully connected Neural Network performs on MNSIT

```
In [0]: # Import MNIST dataset
from keras.datasets import mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
In [170]: print(x_train.shape, y_train.shape)
(60000, 28, 28) (60000,)
```




```

In [171]: # Visualize the data
fig = plt.figure()
for i in range(9):
    plt.subplot(3,3,i+1)
    plt.tight_layout()
    plt.imshow(x_train[i], cmap='gray', interpolation='none')
    plt.title("Digit: {}".format(y_train[i]))
    plt.xticks([])
    plt.yticks([])
fig

```

Out[171]:



```
In [0]: # Flatten the input vectors
x_train = x_train.reshape(60000, 784)
x_test = x_test.reshape(10000, 784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')

# Normalize the input data
x_train /= 255
x_test /= 255
```

```
In [173]: print(x_train.shape, y_train.shape)

(60000, 784) (60000,)
```

```
In [0]: # Now one-hot encode the labels
from keras.utils import np_utils
y_train = np_utils.to_categorical(y_train, 10)
y_test = np_utils.to_categorical(y_test, 10)
```

```
In [175]: print(x_train.shape, y_train.shape)

(60000, 784) (60000, 10)
```

```
In [176]: # Build the NN and train it

model = Sequential()

model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dense(512, activation='relu'))
model.add(Dense(10, activation='softmax'))

model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

h = model.fit(x_train, y_train,
              batch_size=128, epochs=20,
              verbose=2,
              validation_data=(x_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

- 10s - loss: 0.2203 - acc: 0.9349 - val_loss: 0.0998 - val_acc: 0.9699

Epoch 2/20

- 9s - loss: 0.0789 - acc: 0.9756 - val_loss: 0.0892 - val_acc: 0.9699

Epoch 3/20

- 9s - loss: 0.0508 - acc: 0.9836 - val_loss: 0.0705 - val_acc: 0.9778

Epoch 4/20

- 9s - loss: 0.0369 - acc: 0.9880 - val_loss: 0.0704 - val_acc: 0.9792

Epoch 5/20

- 9s - loss: 0.0277 - acc: 0.9909 - val_loss: 0.0746 - val_acc: 0.9770

Epoch 6/20

- 9s - loss: 0.0215 - acc: 0.9926 - val_loss: 0.0861 - val_acc: 0.9759

Epoch 7/20

- 9s - loss: 0.0188 - acc: 0.9941 - val_loss: 0.0773 - val_acc: 0.9817

Epoch 8/20

- 9s - loss: 0.0148 - acc: 0.9950 - val_loss: 0.0813 - val_acc: 0.9796

Epoch 9/20

- 9s - loss: 0.0129 - acc: 0.9958 - val_loss: 0.0874 - val_acc: 0.9799

Epoch 10/20

- 9s - loss: 0.0170 - acc: 0.9945 - val_loss: 0.0919 - val_acc: 0.9773

Epoch 11/20

- 9s - loss: 0.0110 - acc: 0.9966 - val_loss: 0.0877 - val_acc: 0.9778

Epoch 12/20

- 9s - loss: 0.0144 - acc: 0.9953 - val_loss: 0.0858 - val_acc: 0.9813

Epoch 13/20

- 9s - loss: 0.0102 - acc: 0.9966 - val_loss: 0.0861 - val_acc: 0.9806

Epoch 14/20

- 9s - loss: 0.0094 - acc: 0.9967 - val_loss: 0.0926 - val_acc: 0.9790

Epoch 15/20

- 9s - loss: 0.0074 - acc: 0.9974 - val_loss: 0.0908 - val_acc: 0.9815

Epoch 16/20

- 9s - loss: 0.0090 - acc: 0.9972 - val_loss: 0.0962 - val_acc: 0.9807

Epoch 17/20

- 9s - loss: 0.0103 - acc: 0.9966 - val_loss: 0.0900 - val_acc: 0.9811

Epoch 18/20

- 9s - loss: 0.0060 - acc: 0.9977 - val_loss: 0.0949 - val_acc: 0.9818

Epoch 19/20

- 9s - loss: 0.0094 - acc: 0.9970 - val_loss: 0.0900 - val_acc: 0.9839

Epoch 20/20

- 9s - loss: 0.0073 - acc: 0.9979 - val_loss: 0.0925 - val_acc: 0.9809

In [177]: `h.history.keys()`

Out[177]: `dict_keys(['loss', 'acc', 'val_loss', 'val_acc'])`

```
In [179]: # Plot the accuracy during training
plt.plot(h.history['acc'], color='b', label='Accuracy')
plt.plot(h.history['val_acc'], color='r', label='Validation Accuracy')
plt.legend()
```

Out[179]: <matplotlib.legend.Legend at 0x7f108310c710>

